

Final project for the Course High Performance Programming, Philipp Noel von Bachmann

March 16, 2022

1 Introduction

Machine Learning and Data analysis are some of the most influential fields in our current society. One of the main task in Machine Learning is to give a prediction based on some input variables. A common algorithm to use is linear-least-square, which finds the best fit for a linear model. Here, we need an efficient way to calculate this fit. In Data analysis, a common tool is the Principal Component analysis, which tries to reduce the dimensionality of the data in a way, that the reconstruction error gets minimized. PCA relies on an eigenvalue decomposition. For both methods, we can use the QR decomposition too speed up the computation.

2 Problem description

Suppose A is a real, square matrix. Then it can be shown that A can be decomposed into

$$A = QR \tag{1}$$

where Q is orthogonal and R is upper triangular. Finding Q and R is the task of the algorithm.

3 Solution

We will implement an algorithm known as **Givens Rotation**. This algorithm relies on construction a sequences of matrices G_i , such that when multiplying A with G_i , we get a new matrix with a zero at a predefined place. Choosing G_i such that we eliminate the lower diagonal of A , we end up with a R and by multiplying all G_i with Q . We will first show how to eliminate one value at the time by constructing G_i and then how to combine these.

3.1 Givens rotation

First, we define a Givens rotation matrix as

$$G(i, j, \theta) = \begin{bmatrix} 1 & & & & & \\ & \ddots & & & & \\ & & c_{ii} & \cdots & -s_{ij} & \\ & & \vdots & \ddots & \vdots & \\ & & s_{ji} & \cdots & c_{jj} & \\ & & & & & \ddots \\ & & & & & & 1 \end{bmatrix} \quad (2)$$

where $c = \cos(\theta)$, $s = \sin(\theta)$ and any not filled out values are 0. Note: Therefore we can equally represent G by $G(i, j, c, s)$

3.2 Eliminating one value

We will now find θ to solve

$$G(i, j, \theta)^T \begin{bmatrix} \times \\ \vdots \\ a \\ \vdots \\ b \\ \vdots \\ \times \end{bmatrix} = \begin{bmatrix} \times \\ \vdots \\ r \\ \vdots \\ 0 \\ \vdots \\ \times \end{bmatrix} \quad (3)$$

where \times are arbitrary numbers $a, b \in \mathbb{R}$, $r = \sqrt{a^2 + b^2}$. A trivial solution would be

$$c = \frac{a}{r} \quad (4) \quad s = \frac{b}{r} \quad (5)$$

However, r is prone to overflow, so we can instead store it in a different way. If

- $|b| \geq |a|$:
 $t = \frac{a}{b}$, $s = \frac{\text{sgn}(b)}{\sqrt{1+t^2}}$, $c = st$
- else:
 $t = \frac{b}{a}$, $c = \frac{\text{sgn}(a)}{\sqrt{1+t^2}}$, $s = ct$

3.3 Final algorithm

Algorithm 1 Givens rotation

- 1: **procedure** STEP
 - 2: set $R = A$, $Q = I$
 - 3: **for** j in 1 to n **do**
 - 4: **for** i in n down to $j + 1$ **do**
 - 5: compute the Givens rotation $G(i, j, c, s)$ eliminating R_{ij} with $a = i$, $b = i - 1$.
 - 6: set $R = G(i, j, c, s)R$, $Q = QG(i, j, c, s)$
 - 7:
-

Note that we have to do it in this order, since each update affects row and row above.

4 Experiments

4.1 Check for correctness

4.2 Results

4.2.1 Format

In the following section, we will show timing results for different matrix sizes. We denote the sizes by i, j , which we defined as decomposing the matrix $A \in \mathbb{R}^{i \times j}$.

4.2.2 Baseline Results

$i \backslash j$	10	25	50
10	0.0005	0.0008	0.001
25	0.018	0.034	0.049
50	0.235	0.611	1.076

We see that if we vary i , we get a superlinear increase of runtime. If we increase j instead, we see that we get a nearly linear increase in runtime.

(We should only get a minor increase in runtime for j , since we don't need to calculate new 0s). (For i this makes sense, since increasing i by one adds i new 0s to compute.)

5 Random

- we can order them by affecting row and above
- best is probably to create tasks because upper rows require less work
- maybe need memory optimizations, but most of the operations seem to be near anyway

TODO:

- implement checks
- implement good storage of G
- implement efficient matmul of G
- implement one given's rotation
- implement outer loop
- serial optimizations
- parallelizations, probably with openmp

6 Experiments

7 Conclusion

8 References